

Part 1 – MinimaxAgent -- Description

1. Defined a recursive function `minimax()` in `getAction()`.
2. `getAction()` will call `minimax(state, agentIndex=0, cur_depth=0)`

```
agentIndex, cur_depth = 0, 0
_, action = minimax(gameState, agentIndex, cur_depth)
return action
```

3. In `minimax()`:

- ① check if `gameState` is terminal states or not
by `isWin()` or `sLose()` or `(depth == self.depth)`.

-
- ② `next_agentIndex = (agentIndex + 1) % num(agents)`, which is ranging from 0 to `num(agents)-1` since Pacman has `agentIndex = 0` and Ghosts have `agentIndex > 0`.

- ③ `next_depth` is increased based on rounds.

`next_depth += 1` at Pacman's turns.

```
#----- terminal state -----
if gameState.isWin() or gameState.isLose() or cur_depth == self.depth :
    return self.evaluationFunction(gameState), None
#----- pre-defined -----
next_agentIndex = (agentIndex + 1) % gameState.getNumAgents()
next_depth = (cur_depth + 1) if (next_agentIndex == 0) else cur_depth
```

-
- ④ If `agentIndex == 0` (Pacman's turn):

Loop through all legal actions and `return (max_value, max_action)`

Else (Ghost's turn):

Loop through all legal actions and `return (min_value, min_action)`

- ⑤ What we do for each action in the above loops is
getting `next_state` by `gameState.getNextState()` and then getting
`next_state's` value by calling `minimax()`.

```
#----- Pacman -----
if(agentIndex==0):
    # default value, action
    max_value = float("-inf")
    max_action = None
    # Loop through legal actions to get the maximal value
    # For each action, we can get a next_state and next_state's value
    # If next_state's value is the larger than other past states's value, save it
    for action in gameState.getLegalActions(agentIndex):
        next_state = gameState.getNextState(agentIndex, action)
        next_value, _ = minimax(next_state, next_agentIndex, cur_depth)
        if next_value > max_value:
            max_value = next_value
            max_action = action
    return max_value, max_action
#----- Ghost -----
else:
    min_value = float("inf")
    min_action = None
    for action in gameState.getLegalActions(agentIndex):
        next_state = gameState.getNextState(agentIndex, action)
        next_value, _ = minimax(next_state, next_agentIndex, next_depth)
        if next_value < min_value:
            min_value = next_value
            min_action = action
    return min_value, min_action
```

Part 1 – MinimaxAgent -- Code

```
class MinimaxAgent(MultiAgentSearchAgent):

    def getAction(self, gameState):
        # Begin your code (Part 1)
        def minmax(gameState, agentIndex, cur_depth):
            #----- terminal state -----
            if gameState.isWin() or gameState.isLose() or cur_depth == self.depth :
                return self.evaluationFunction(gameState), None
            #----- pre-defined -----
            next_agentIndex = (agentIndex + 1) % gameState.getNumAgents()
            next_depth = (cur_depth + 1) if (next_agentIndex == 0) else cur_depth
            #----- Pacman -----
            if(agentIndex==0):
                # default value, action
                max_value = float("-inf")
                max_action = None
                # Loop through legal actions to get the maximal value
                # For each action, we can get a next_state and next_state's value
                # If next_state's value is the larger than other past states's value, save it
                for action in gameState.getLegalActions(agentIndex):
                    next_state = gameState.getNextState(agentIndex, action)
                    next_value, _ = minmax(next_state, next_agentIndex, cur_depth)
                    if next_value > max_value:
                        max_value = next_value
                        max_action = action
                return max_value, max_action
            #----- Ghost -----
            else:
                min_value = float("inf")
                min_action = None
                for action in gameState.getLegalActions(agentIndex):
                    next_state = gameState.getNextState(agentIndex, action)
                    next_value, _ = minmax(next_state, next_agentIndex, next_depth)
                    if next_value < min_value:
                        min_value = next_value
                        min_action = action
                return min_value, min_action
            #-----
            agentIndex, cur_depth = 0, 0
        _, action = minmax(gameState, agentIndex, cur_depth)
        return action
        # End your code (Part 1)
```

Part2 – AlphaBetaAgent -- Description

1. There are 3 functions in class AlphaBetaAgent. They are `getAction()`, `max_value()`, and `min_value()`.
2. `getAction()` call `max_value(state, index=0, depth=0, alpha=-inf, Beta=inf)`

```
def getAction(self, gameState):
    agentIndex = 0
    cur_depth = 0
    _, action = self.max_value(gameState, agentIndex, cur_depth, float('-inf'), float('inf'))
    return action
```

3. In `max_value()` :

- ① Check terminal states.

`next_agentIndex` and `next_depth` are mentioned in Part0.

```
if gameState.isWin() or gameState.isLose() or cur_depth == self.depth :
    return self.evaluationFunction(gameState), None

next_agentIndex = (agentIndex + 1) % gameState.getNumAgents()
next_depth = (cur_depth + 1) if (next_agentIndex == 0) else cur_depth
```

- ② Loop through all legal actions of current state. For each action:

- (1) get `next_state` by calling `getNextState()`
- (2) get `next_state`'s value by calling `max_value()` if `next_agentIndex == 0` ; otherwise call `min_value()`.
- (3) After `next_state`'s value being returned, update `max_value`, `alpha`.
- (4) If (`max_value > Beta`) : return immediately to prune the state tree.

```
# Loop through actions to get the max_value
max_value = float("-inf")
max_action = None
for action in gameState.getLegalActions(agentIndex):
    (1) next_state = gameState.getNextState(agentIndex, action)
    # if next_state is Pacman's turn, call max_value() to get next_state's value
    (2) if (next_agentIndex==0):
        next_value, _ = self.max_value(next_state, next_agentIndex, next_depth, alpha, beta)
    # if next_state is Ghost's turn, call min_value() to get next_state's value
    else:
        next_value, _ = self.min_value(next_state, next_agentIndex, next_depth, alpha, beta)
    (3) # update max_value
    if next_value > max_value:
        max_value = next_value
        max_action = action
    # update alpha
    alpha = max(alpha, max_value)
    (4) # Pruning: if max_value > Beta
    if max_value > beta:
        return max_value, action
return max_value, max_action
```

4. In `min_value()` :

- ①② steps are same as `max_value()`.
- ③ Loop through all legal actions of current state. For each action:
 - (1)(2) are same as `max_value()`.
 - (3) After `next_state`'s value being returned, update `min_value`, `beta`.
 - (4) If (`min_value < alpha`) : return immediately to prune the state tree.

```
# update min_value
if next_value < min_value:
    min_value = next_value
    min_action = action
# update beta
beta = min(beta, min_value)
# Pruning: if min_value < alpha
if min_value < alpha:
    return min_value, action
return min_value, min_action
```

Part 2 – AlphaBetaAgent -- Code

```
class AlphaBetaAgent(MultiAgentSearchAgent):
    """
    Your minimax agent with alpha-beta pruning
    """
    def getAction(self, gameState):
        agentIndex = 0
        cur_depth = 0
        _, action = self.max_value(gameState, agentIndex, cur_depth, float('-inf'), float('inf'))
        return action

    def max_value(self, gameState, agentIndex, cur_depth, alpha, beta):
        # terminal states
        if gameState.isWin() or gameState.isLose() or cur_depth == self.depth:
            return self.evaluationFunction(gameState), None
        # pre-defined next_state's info
        next_agentIndex = (agentIndex + 1) % gameState.getNumAgents()
        next_depth = (cur_depth + 1) if (next_agentIndex == 0) else cur_depth
        # loop through actions to get the max_value
        max_value = float("-inf")
        max_action = None
        for action in gameState.getLegalActions(agentIndex):
            next_state = gameState.getNextState(agentIndex, action)
            # if next_state is Pacman's turn, call max_value() to get next_state's value
            if (next_agentIndex == 0):
                next_value, _ = self.max_value(next_state, next_agentIndex, next_depth, alpha, beta)
            # if next_state is Ghost's turn, call min_value() to get next_state's value
            else:
                next_value, _ = self.min_value(next_state, next_agentIndex, next_depth, alpha, beta)
            # update max_value
            if next_value > max_value:
                max_value = next_value
                max_action = action
            # Pruning: if max_value > Beta
            alpha = max(alpha, max_value)
            if max_value > beta:
                return max_value, action
        return max_value, max_action

    def min_value(self, gameState, agentIndex, cur_depth, alpha, beta):
        if gameState.isWin() or gameState.isLose() or cur_depth == self.depth:
            return self.evaluationFunction(gameState), None

        next_agentIndex = (agentIndex + 1) % gameState.getNumAgents()
        next_depth = (cur_depth + 1) if (next_agentIndex == 0) else cur_depth
        # loop through actions to get the min_value
        min_value = float("inf")
        min_action = None
        for action in gameState.getLegalActions(agentIndex):
            next_state = gameState.getNextState(agentIndex, action)
            # if next_state is Pacman's turn, call max_value() to get next_state's value
            if (next_agentIndex == 0):
                next_value, _ = self.max_value(next_state, next_agentIndex, next_depth, alpha, beta)
            # if next_state is Ghost's turn, call min_value() to get next_state's value
            else:
                next_value, _ = self.min_value(next_state, next_agentIndex, next_depth, alpha, beta)
            # update min_value
            if next_value < min_value:
                min_value = next_value
                min_action = action
            # Pruning: if min_value < alpha
            beta = min(beta, min_value)
            if min_value < alpha:
                return min_value, action
        return min_value, min_action
```

Part3 – ExpectimaxAgent -- Description

1. Define a recursive function `expectimax()` in `getAction()`.
2. `getAction()` will call `expectimax(state, agentIndex=0, cur_depth=0)`

```
agentIndex = 0
cur_depth = 0
_, action = self.expectimax(gameState, agentIndex, cur_depth)
return action
```

3. In `expectimax()` :

- ① check `terminal states`, predefine `next_agentIndex` and `next_depth` as `MinimaxAgent` and `AlphaBeataAgent`.

```
#----- terminal state -----
if game_state.isWin() or game_state.isLose() or depth == self.depth :
    return self.evaluationFunction(game_state), None
#----- pre-defined -----
next_agentIndex = (agentIndex + 1) % game_state.getNumAgents()
next_depth = (depth + 1) if (next_agentIndex == 0) else depth
LegalActions = game_state.getLegalActions(agentIndex)
```

-
- ② If `agentIndex == 0` (Pacman's turn) :

Loop through all legal actions and `return (max_value, max_action)`

Else (Ghost's turn) :

Loop through all legal actions and `return (expected_value, expected_action)`

- ③ The loop in Pacman's part is same as what applied in `MinimaxAgent`'s.
- ④ The loop in Ghost's part `return (expected_value, expected_action)`, where $\text{expected value} = \sum_{i=1}^{\text{num}(\text{next_states})} p * \text{value}_i$ with $p = 1/\text{num}(\text{next_states})$.

```
#----- Pacman -----
if agentIndex == 0:
    max_value = float("-inf")
    max_action = None

    for action in LegalActions:
        next_state = game_state.getNextState(agentIndex, action)
        next_value, _ = self.expectimax(next_state, next_agentIndex, next_depth)
        if next_value > max_value:
            max_value = next_value
            max_action = action
    return max_value, max_action
#----- Ghost -----
else:
    expected_value = 0
    expected_action = None
    probability = 1.0 / len(LegalActions)

    for action in LegalActions:
        next_state = game_state.getNextState(agentIndex, action)
        next_value, _ = self.expectimax(next_state, next_agentIndex, next_depth)
        expected_value += probability * next_value
    return expected_value, expected_action
```

Part3 – ExpectimaxAgent -- Code

```
class ExpectimaxAgent(MultiAgentSearchAgent):
    """
    Your expectimax agent
    """
    def getAction(self, gameState):
        agentIndex = 0
        cur_depth = 0
        _, action = self.expectimax(gameState, agentIndex, cur_depth)
        return action

    def expectimax(self, game_state, agentIndex, depth):
        #----- terminal state -----
        if game_state.isWin() or game_state.isLose() or depth == self.depth :
            return self.evaluationFunction(game_state), None
        #----- pre-defined -----
        next_agentIndex = (agentIndex + 1) % game_state.getNumAgents()
        next_depth = (depth + 1) if (next_agentIndex == 0) else depth
        LegalActions = game_state.getLegalActions(agentIndex)
        #----- Pacman -----
        if agentIndex == 0:
            max_value = float("-inf")
            max_action = None

            for action in LegalActions:
                next_state = game_state.getNextState(agentIndex, action)
                next_value, _ = self.expectimax(next_state, next_agentIndex, next_depth)
                if next_value > max_value:
                    max_value = next_value
                    max_action = action
            return max_value, max_action
        #----- Ghost -----
        else:
            expected_value = 0
            expected_action = None
            probability = 1.0 / len(LegalActions)

            for action in LegalActions:
                next_state = game_state.getNextState(agentIndex, action)
                next_value, _ = self.expectimax(next_state, next_agentIndex, next_depth)
                expected_value += probability * next_value
            return expected_value, expected_action
```

Part4 – betterEvaluationFunction -- Description & Code

1. Assume 2 kinds of factors : `min_food_dist` and `game_score`
2. `factor1` -- `min_food_dist` :
 - ① look for the minimal distance from pacman to food.
 - ② pacman's position is retrieved by `getPacmanPosition()`, and food's position are retrieved by `getFood()`.
 - ③ The function used for calculating distance between Pacman and Food is `manhattanDistance()`, which is imported at default.
3. `factor2` -- `game_score` is retrieved by `getScore()`
4. The weight for factors is set by try and error.
5. Since smaller `min_food_dist` is better, `factor1` is set by `1.0/(1.0+min_food_dist)`. Add `min_food_dist` by 1 to avoid division by zero.
6. The final evaluated value = `factor1*weight1 + factor2*weight2`.

```
def betterEvaluationFunction(currentGameState):
    """
    Your extreme evaluation function
    """
    # 2 kinds of possible factors
    min_food_dist = float('inf')
    game_score = currentGameState.getScore()

    # the first factor -- min_food_dist
    pacman_pos = currentGameState.getPacmanPosition()
    for food_position in currentGameState.getFood().asList():
        pac_food_dist = manhattanDistance(pacman_pos, food_position)
        if(pac_food_dist < min_food_dist):
            min_food_dist = pac_food_dist

    # modify weights manually
    min_food_dist_weight = 1
    game_score_weight = 1

    # Final_value = Sum(factor_i * weight_i)
    Final_value = 1.0/(1.0+min_food_dist) * min_food_dist_weight \
        + game_score * game_score_weight

    return Final_value
```

Part5 – problems you meet and how you solve them

The biggest problem is the definition of provided function.

Though there's a short usage description at the top of MinimaxAgent code, what objects those function will return are still confusing. I believe we should go to `pacman.py` and `utils.py` frequently to look for the answers. By the way, the type hint is very helpful in exploring default function and class members(X

~~Praise for type hint~~